

Introduction to R for Biologists

Maria Doyle, Jessica Chung, Vicky Perreau

14 July 2023

Contents

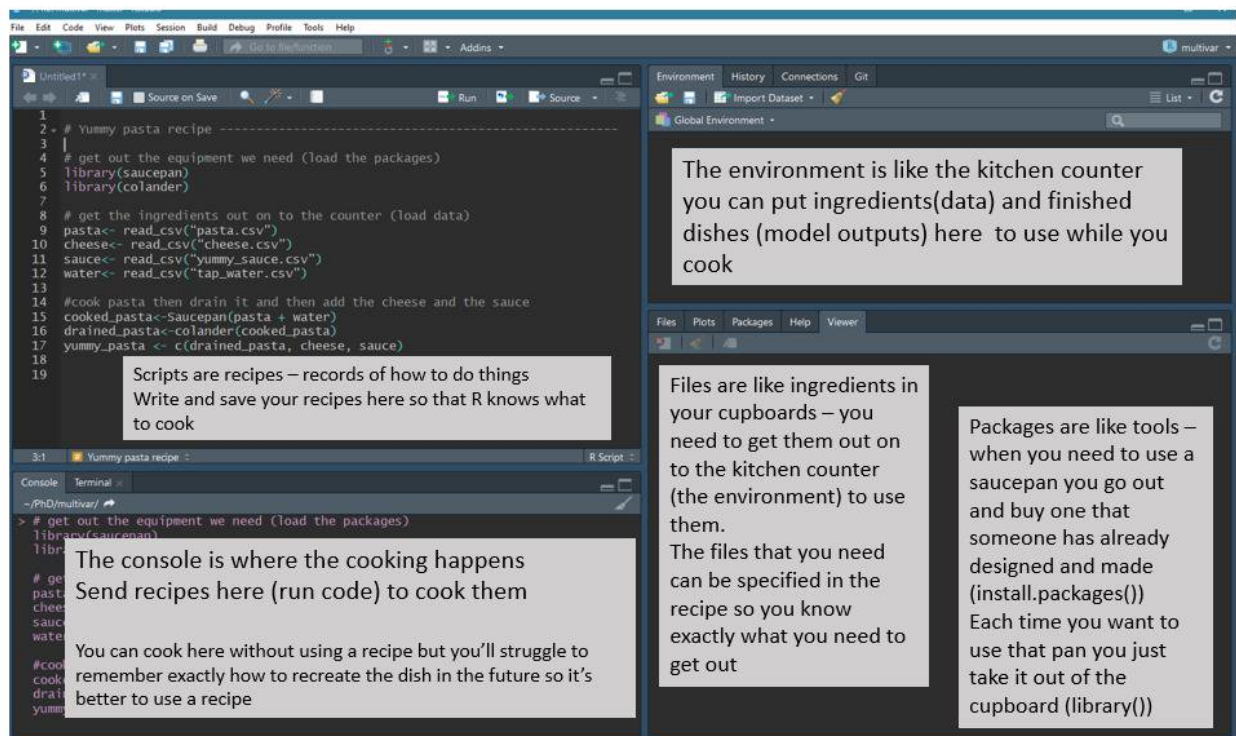
R for Biologists course	2
Intro to R and RStudio	2
R script vs console	2
Working directory	3
Packages	4
Getting help	4
Common R errors	4
Getting started with data	5
Data files	5
GREIN (GEO RNA-seq Experiments Interactive Navigator)	5
RNA-seq dataset	5
Tidyverse	6
Loading the data	7
Getting to know the data	9
Formatting the data	12
Converting from wide to long format	12
Joining two tables	14
Plotting with ggplot2	16
Creating a boxplot	16
Colouring by categories	18
Creating subplots for each gene	19
Make shorter category names	19
Filter for genes of interest	20
Create plots for each gene	20
Customising the plot	22
Specifying colours	22
Axis labels and Title	23
Themes	24
Order of categories	26
Saving plots	28
Session Info	28

R for Biologists course

R takes time to learn, like a spoken language. No one can expect to be an R expert after learning R for a few hours. This course has been designed to introduce biologists to R, showing some basics, and also some powerful things R can do (things that would be more difficult to do with Excel). The aim is to give beginners the confidence to continue learning R, so the focus here is on tidyverse and visualisation of biological data, as we believe this is a productive and engaging way to start learning R. After this short introduction you could use this book to dive a bit deeper.

Intro to R and RStudio

RStudio is an interface that makes it easier to use R. There are four windows in RStudio. The screenshot below shows an analogy linking the different RStudio windows to cooking.



R script vs console

There are two ways to work in RStudio in the console or in a script. We can type a command in the console and press **Enter** to run it. Try running the command below in the console.

```
1 + 1
```

```
## [1] 2
```

Or we can use an R script. To create a script, from the top menu in RStudio: **File > New File > R Script**. Now type the command below in the script. This time, to run the command, you use **Ctrl + Enter** for

Windows/Linux or **Cmd + Enter** for MacOS. This sends the command where the cursor is from the script to the console. You can highlight multiple commands and then press **Cmd/Ctrl + Enter** to run them one after the other.

```
2 + 2
```

```
## [1] 4
```

As the RStudio screenshot above explains, if we work in the console we don't have a good record (recipe) of what we've done. We can see commands we've run in the History panel (top right window), and we can go backwards and forwards through our history in the console using the up arrow and down arrow. But the history includes everything we've tried to run, including our mistakes so it is good practice to use an R script.

We can also add comments to a script. These are notes to yourself or others about the commands in the script. Comments start with a **#** which tells R not to run them as commands.

```
# testing R
```

```
2 + 2
```

```
## [1] 4
```


Keeping an accurate record of how you've manipulated your data is important for reproducible research. Writing detailed comments and documenting your work are useful reminders to your future self (and anyone else reading your scripts) on what your code does.

Working directory

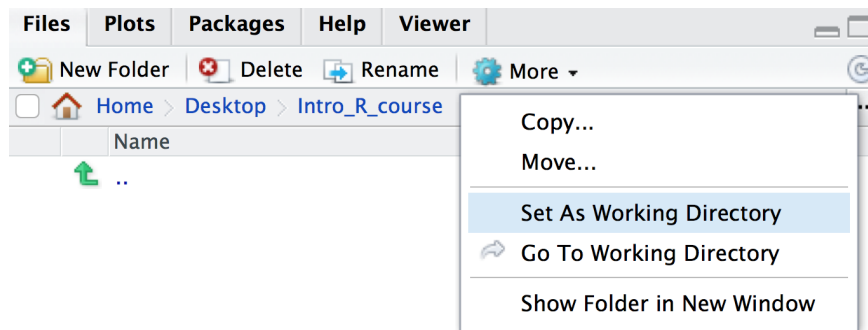
Opening an RStudio session launches it from a specific location. This is the 'working directory'. **R looks in the working directory by default to read in data and save files.** You can find out what the working directory is by using the command `getwd()`. This shows you the path to your working directory in the console. In Mac this is in the format `/path/to/working/directory` and in Windows `C:\path\to\working\directory`. It is often useful to have your data and R scripts in the same directory and set this as your working directory. We will do this now.

Make a folder for this course somewhere on your computer that you will be able to easily find. Name the folder for example, `Intro_R_course`. Then, to set this folder as your working directory:

In RStudio click on the 'Files' tab and then click on the three dots, as shown below.



In the window that appears, find the folder you created (e.g. `Intro_R_course`), click on it, then click 'Open'. The files tab will now show the contents of your new folder. Click on **More > Set As Working Directory**, as shown below.



Save the script you created in the previous section as `intro.R` in this directory. You can do this by clicking on **File > Save** and the default location should be the current working directory (e.g. `Intro_R_course`).

Note: You can use an RStudio project as described here to automatically keep track of and set the working directory.

Packages

If it's not already installed on your computer, you can use the `install.packages` function to install a **package**. A package is a collection of functions along with documentation, code, tests and example data.

```
install.packages("tidyverse")
```

We will see many functions in this tutorial. Functions are “canned scripts” that automate more complicated sets of commands. Many functions are predefined, or can be made available by importing R packages. A function usually takes one or more inputs called *arguments*. Here `tidyverse` is the argument to the `install.packages()` function.

Note: functions require parentheses after the function name.

Getting help

To see what any function in R does, type a `?` before the name and help information will appear in the Help panel on the right in RStudio. Or you can search the function name in the Help panel search box. Google and Stack Overflow are also useful resources for getting help.

```
?install.packages
```

[INFO] Tab completion

A very useful feature is Tab completion. You can start typing and use Tab to autocomplete code, for example, a function name.

Common R errors

R error messages are common and can sometimes be cryptic. You most likely will encounter at least one error message during this tutorial. Some common reasons for errors are:

- Case sensitivity. In R, as in other programming languages, case sensitivity is important. `?install.packages` is different to `?Install.packages`.
- Missing commas
- Mismatched parentheses or brackets
- Not quoting file paths

- Not finishing a command so seeing “+” in the console. If you need to, you can press ESC to cancel the command.

To see examples of some R error messages with explanations see [here](#)

Getting started with data

Data files

The data files required for this workshop are available on GitHub. To download the data.zip file, you can [click here](#). Unzip the file and store this `data` folder in your working directory.


GREIN (GEO RNA-seq Experiments Interactive Navigator)

In this tutorial, we will learn some R through creating plots to visualise data from an RNA-seq experiment. RNA-seq counts file can be obtained from the GREIN platform. GREIN provides >6,500 published datasets from GEO that have been uniformly processed. It is available at <http://www.ilincs.org/apps/grein/>. You can search for a dataset of interest using the GEO code. We obtained the dataset used here using the code GSE60450. GREIN provide QC metrics for the RNA-seq datasets and both raw and normalized counts. We will use the normalized counts here. These are the counts of reads for each gene for each sample normalized for differences in sequencing depth and composition bias. Generally, the higher the number of counts the more the gene is expressed.

RNA-seq dataset

Here we will create some plots using RNA-seq data from the paper by Fu et al. 2015, GEO code GSE60450. This study examined expression in basal and luminal cells from mice at different stages (virgin, pregnant and lactating). There are 2 samples per group and 6 groups, 12 samples in total.

RNA-seq of Mouse mammary gland




Fu *et al.* (2015) 'EGF-mediated induction of Mcl-1 at the switch to lactation is essential for alveolar cell survival' Nat Cell Biol

Tidyverse



In this course we will use the **tidyverse**. The tidyverse is a collection of R packages that includes the extremely widely used **ggplot2** package.



R packages for data science

The tidyverse is an opinionated [collection of R packages](#) designed for data science. All packages share an underlying design philosophy, grammar, and data structures.

The tidyverse makes data science faster, easier and more fun.

Why tidyverse? Why tidy data? Why is it such a game-changer?

Loading the data

We use `library()` to load in the packages that we need. As described in the cooking analogy in the first screenshot, `install.packages()` is like buying a saucepan, `library()` is taking it out of the cupboard to use it.

```
library(tidyverse)
```

The files we will use are csv comma-separated, so we will use the `read_csv()` function from the tidyverse. There is also a `read_tsv()` function for tab-separated values.

We will use the counts file called `GSE60450_GeneLevel_Normalized(CPM.and.TMM)_data.csv` that's in a folder called `data` i.e. the path to the file should be `data/GSE60450_GeneLevel_Normalized(CPM.and.TMM)_data.csv`.

We can read the counts file into R with the command below. We'll store the contents of the counts file in an **object** called `counts`. This stores the file contents in R's memory making it easier to use.

```
# read in counts file
counts <- read_csv("data/GSE60450_GeneLevel_Normalized(CPM.and.TMM)_data.csv")

## New names:
## Rows: 23735 Columns: 14
## -- Column specification
## ----- Delimiter: "," chr
## (2): ...1, gene_symbol dbl (12): GSM1480291, GSM1480292, GSM1480293,
## GSM1480294, GSM1480295, GSM148...
## i Use `spec()` to retrieve the full column specification for this data. i
## Specify the column types or set `show_col_types = FALSE` to quiet this message.
## * `` -> `...1`

# read in metadata
sampleinfo <- read_csv("data/GSE60450_filtered_metadata.csv")
```

```
## New names:
## Rows: 12 Columns: 4
## -- Column specification
## ----- Delimiter: "," chr
## (4): ...1, characteristics, immunophenotype, developmental stage
## i Use `spec()` to retrieve the full column specification for this data. i
## Specify the column types or set `show_col_types = FALSE` to quiet this message.
## * `` -> `...1`
```

There is some information output by `read_csv` on “column specification”. It tells us that there is a missing column name in the header and it has been filled with the name “...1”, which is how `read_csv` handles missing column names by default. It also tells us what data types `read_csv` is detecting in each column. Columns with text characters have been detected (`col_character`) and also columns with numbers (`col_double`). We won’t get into the details of R data types in this tutorial but they are important to know and you can read more about them in the R for Data Science book.

In R we use `<-` to assign values to objects. `<-` is the **assignment operator**. It assigns values on the right to objects on the left. So to create an object, we need to give it a name (e.g. `counts`), followed by the assignment operator `<-`, and the value we want to give it. We can give an object almost any name we want but there are some rules and conventions as described in the tidyverse R style guide

We can read in a file from a path on our computer or on the web and use this as the value. Note that we need to put quotes (“”) around file paths.

[INFO] Assignment operator shortcut

In RStudio, typing `Alt + -` (holding down `Alt` at the same time as the `-` key) will write `<-` in a single keystroke in Windows, while typing `> Option + -` (holding down `Option` at the same time as the `-` key) does the same in a Mac.

Exercise

1. Test what happens if you type `Library(tidyverse)`
What is wrong and how would you fix it?
2. Test what happens if you type `library(tidyverse)`
What is wrong and how would you fix it?
3. Test what happens if you type `library(tidyverse)`
What is wrong and how would you fix it?
4. Test what happens if you type
`read_tsv("data/GSE60450_filtered_metadata.csv")`
What is wrong and how would you fix it?
5. Test what happens if you type
`read_csv("data/GSE60450_filtered_metadata.csv")`
What is wrong and how would you fix it?
6. Test what happens if you type
`read_csv("GSE60450_filtered_metadata.csv")`
What is wrong and how would you fix it?
7. What is the name of the first column you get with each of these 2 commands?
`read.csv("data/GSE60450_filtered_metadata.csv")`


```
and
read_csv("data/GSE60450_filtered_metadata.csv")
```

8. If you run

```
read_csv("data/GSE60450_filtered_metadata.csv")
```

what is the difference between the column header you see `developmental stage` and `'developmental stage'`?

Getting to know the data

When assigning a value to an object, R does not print the value. For example, here we don't see what's in the counts or sampleinfo files. But there are ways we can look at the data. We will demonstrate using the `sampleinfo` object.

We can type the name of the object and this will print the first few lines and some information, such as number of rows.

```
sampleinfo
```

```
## # A tibble: 12 x 4
##   ...1      characteristics      immunoph~1 devel~2
##   <chr>      <chr>              <chr>      <chr>
## 1 GSM1480291 mammary gland, luminal cells, virgin      luminal ~ virgin
## 2 GSM1480292 mammary gland, luminal cells, virgin      luminal ~ virgin
## 3 GSM1480293 mammary gland, luminal cells, 18.5 day pregnancy luminal ~ 18.5 d~
## 4 GSM1480294 mammary gland, luminal cells, 18.5 day pregnancy luminal ~ 18.5 d~
## 5 GSM1480295 mammary gland, luminal cells, 2 day lactation luminal ~ 2 day ~
## 6 GSM1480296 mammary gland, luminal cells, 2 day lactation luminal ~ 2 day ~
## 7 GSM1480297 mammary gland, basal cells, virgin      basal ce~ virgin
## 8 GSM1480298 mammary gland, basal cells, virgin      basal ce~ virgin
## 9 GSM1480299 mammary gland, basal cells, 18.5 day pregnancy basal ce~ 18.5 d~
## 10 GSM1480300 mammary gland, basal cells, 18.5 day pregnancy basal ce~ 18.5 d~
## 11 GSM1480301 mammary gland, basal cells, 2 day lactation basal ce~ 2 day ~
## 12 GSM1480302 mammary gland, basal cells, 2 day lactation basal ce~ 2 day ~
## # ... with abbreviated variable names 1: immunophenotype,
## # 2: `developmental stage`
```

We can also use `dim()` to see the dimensions of an object, the number of rows and columns.

```
dim(sampleinfo)
```

```
## [1] 12 4
```

This show us there are 12 rows and 4 columns.

In the Environment Tab in the top right panel in RStudio we can also see the number of rows and columns in the objects we have in our session.

We can also take a look the first few lines with `head()`. This shows us the first 6 lines.

```
head(sampleinfo)
```

```
## # A tibble: 6 x 4
##   ...1      characteristics      immunoph~1 devel~2
##   <chr>      <chr>              <chr>      <chr>
## 1 GSM1480291 mammary gland, luminal cells, virgin      luminal c~ virgin
## 2 GSM1480292 mammary gland, luminal cells, virgin      luminal c~ virgin
## 3 GSM1480293 mammary gland, luminal cells, 18.5 day pregnancy luminal c~ 18.5 d~
## 4 GSM1480294 mammary gland, luminal cells, 18.5 day pregnancy luminal c~ 18.5 d~
```

```
## 5 GSM1480295 mammary gland, luminal cells, 2 day lactation luminal c~ 2 day ~
## 6 GSM1480296 mammary gland, luminal cells, 2 day lactation luminal c~ 2 day ~
## # ... with abbreviated variable names 1: immunophenotype,
## # 2: `developmental stage`
```

We can look at the last few lines with `tail()`. This shows us the last 6 lines. This can be useful to check the bottom of the file, that it looks ok.

```
tail(sampleinfo)
```

```
## # A tibble: 6 x 4
##   ...1      characteristics immunophen~1 devel~2
##   <chr>      <chr>          <chr>      <chr>
## 1 GSM1480297 mammary gland, basal cells, virgin basal cell ~ virgin
## 2 GSM1480298 mammary gland, basal cells, virgin basal cell ~ virgin
## 3 GSM1480299 mammary gland, basal cells, 18.5 day pregnancy basal cell ~ 18.5 d~
## 4 GSM1480300 mammary gland, basal cells, 18.5 day pregnancy basal cell ~ 18.5 d~
## 5 GSM1480301 mammary gland, basal cells, 2 day lactation basal cell ~ 2 day ~
## 6 GSM1480302 mammary gland, basal cells, 2 day lactation basal cell ~ 2 day ~
## # ... with abbreviated variable names 1: immunophenotype,
## # 2: `developmental stage`
```

Or we can see the whole file with `View()`.

```
View(sampleinfo)
```

In the Environment tab we can see how many rows and columns the object contains and we can click on the icon to view all the contents in a tab. This runs the command `View()` for us.

We can see all the column names with `colnames()`.

```
colnames(sampleinfo)
```

```
## [1] "...1"      "characteristics"      "immunophenotype"
## [4] "developmental stage"
```

We can access individual columns by name using the `$` symbol. For example we can see what's contained in the characteristics column.

```
sampleinfo$characteristics
```

```
## [1] "mammary gland, luminal cells, virgin"
## [2] "mammary gland, luminal cells, virgin"
## [3] "mammary gland, luminal cells, 18.5 day pregnancy"
## [4] "mammary gland, luminal cells, 18.5 day pregnancy"
## [5] "mammary gland, luminal cells, 2 day lactation"
## [6] "mammary gland, luminal cells, 2 day lactation"
## [7] "mammary gland, basal cells, virgin"
## [8] "mammary gland, basal cells, virgin"
## [9] "mammary gland, basal cells, 18.5 day pregnancy"
## [10] "mammary gland, basal cells, 18.5 day pregnancy"
## [11] "mammary gland, basal cells, 2 day lactation"
## [12] "mammary gland, basal cells, 2 day lactation"
```

If we just wanted to see the first 3 values in the column we can specify this using square brackets. Obtaining a selection of values this way is called 'subsetting'.

```
sampleinfo$characteristics[1:3]
```

```
## [1] "mammary gland, luminal cells, virgin"
```

```
## [2] "mammary gland, luminal cells, virgin"
## [3] "mammary gland, luminal cells, 18.5 day pregnancy"
```

In the previous section, when we loaded in the data from the csv file, we noticed that the first column had a missing column name and by default, `read_csv` function assigned a name of “...1” to it. Let’s change this column to something more descriptive now. We can do this by combining a few things we’ve just learnt.

First, we use the `colnames()` function to obtain the column names of `sampleinfo`. Then we use square brackets to subset the first value of the column names (`[1]`). Last, we use the assignment operator (`<-`) to set the new value of the first column name to “`sample_id`”.

```
colnames(sampleinfo)[1] <- "sample_id"
```

Let’s check if this has been changed correctly.

```
sampleinfo

## # A tibble: 12 x 4
##   sample_id characteristics immunop~1 devel~2
##   <chr>      <chr>          <chr>      <chr>
## 1 GSM1480291 mammary gland, luminal cells, virgin luminal ~ virgin
## 2 GSM1480292 mammary gland, luminal cells, virgin luminal ~ virgin
## 3 GSM1480293 mammary gland, luminal cells, 18.5 day pregnancy luminal ~ 18.5 d~
## 4 GSM1480294 mammary gland, luminal cells, 18.5 day pregnancy luminal ~ 18.5 d~
## 5 GSM1480295 mammary gland, luminal cells, 2 day lactation luminal ~ 2 day ~
## 6 GSM1480296 mammary gland, luminal cells, 2 day lactation luminal ~ 2 day ~
## 7 GSM1480297 mammary gland, basal cells, virgin basal ce~ virgin
## 8 GSM1480298 mammary gland, basal cells, virgin basal ce~ virgin
## 9 GSM1480299 mammary gland, basal cells, 18.5 day pregnancy basal ce~ 18.5 d~
## 10 GSM1480300 mammary gland, basal cells, 18.5 day pregnancy basal ce~ 18.5 d~
## 11 GSM1480301 mammary gland, basal cells, 2 day lactation basal ce~ 2 day ~
## 12 GSM1480302 mammary gland, basal cells, 2 day lactation basal ce~ 2 day ~
## # ... with abbreviated variable names 1: immunophenotype,
## # 2: `developmental stage`
```

The first column is now named “`sample_id`”.

We can also do the same to the counts data. This time, we rename the first column name from “...1” to “`gene_id`”.

```
colnames(counts)[1] <- "gene_id"
```

[INFO] Multiple methods

There are multiple ways to rename columns. We’ve covered one way here, but another way is using the `rename()` function. When programming, you’ll often find many ways to do the same thing. Often there is one obvious method depending on the context you’re in.

Other useful commands for checking data are `str()` and `summary()`.

`str()` shows us the structure of our data. It shows us what columns there are, the first few entries, and what data type they are e.g. character or numbers (double or integer).

```
str(sampleinfo)

## spc_tbl_ [12 x 4] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ sample_id      : chr [1:12] "GSM1480291" "GSM1480292" "GSM1480293" "GSM1480294" ...
## $ characteristics : chr [1:12] "mammary gland, luminal cells, virgin" "mammary gland, luminal cells, 18.5 day pregnancy" ...
```

```
## $ immunophenotype      : chr [1:12] "luminal cell population" "luminal cell population" "luminal cell
## $ developmental stage: chr [1:12] "virgin" "virgin" "18.5 day pregnancy" "18.5 day pregnancy" ...
## - attr(*, "spec")=
## .. cols(
## ..   ...1 = col_character(),
## ..   characteristics = col_character(),
## ..   immunophenotype = col_character(),
## ..   `developmental stage` = col_character()
## .. )
## - attr(*, "problems")=<externalptr>
```

`summary()` generates summary statistics of our data. For numeric columns (columns of type double or integer) it outputs statistics such as the min, max, mean and median. We will demonstrate this with the counts file as it contains numeric data. For character columns it shows us the length (how many rows).

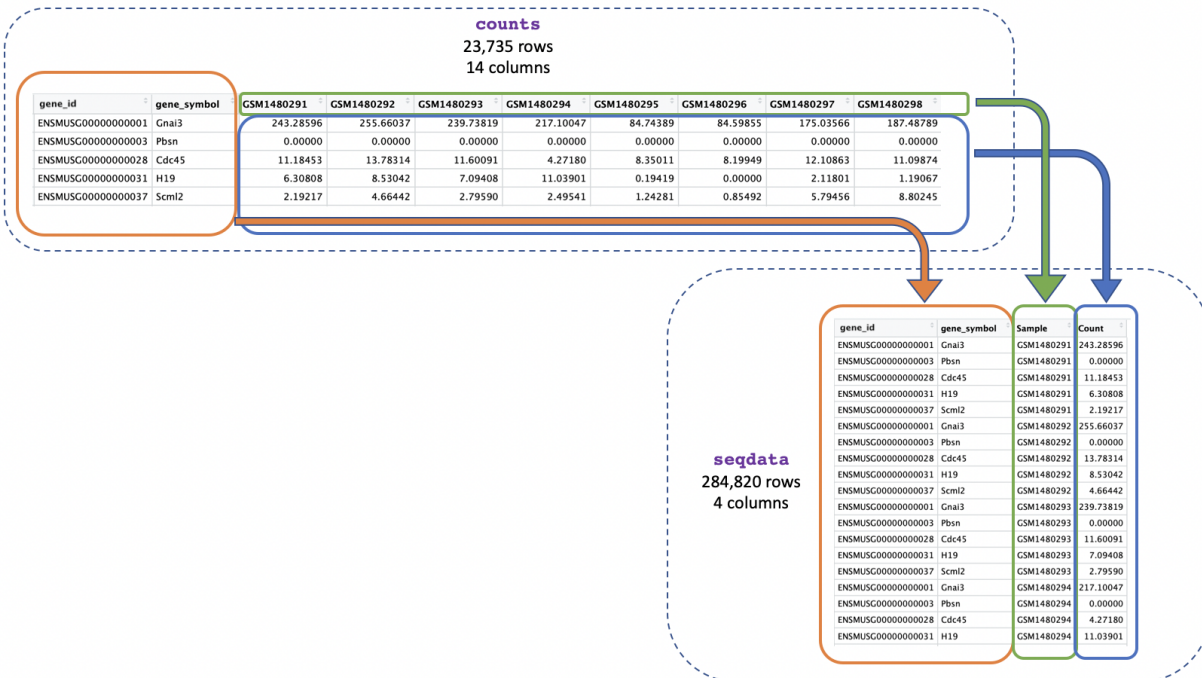
```
summary(counts)
```

```
##      gene_id      gene_symbol      GSM1480291      GSM1480292
## Length:23735      Length:23735      Min.   :    0.000      Min.   :    0.000
## Class :character      Class :character      1st Qu.:    0.000      1st Qu.:    0.000
## Mode  :character      Mode  :character      Median :    1.745      Median :    1.891
##                                     Mean  :   42.132      Mean   :   42.132
##                                     3rd Qu.:   29.840      3rd Qu.:   29.604
##                                     Max.   :12525.066      Max.    :12416.211
##      GSM1480293      GSM1480294      GSM1480295      GSM1480296
## Min.   :    0.00      Min.   :    0.00      Min.   :    0.00      Min.   :    0.00
## 1st Qu.:    0.00      1st Qu.:    0.00      1st Qu.:    0.00      1st Qu.:    0.00
## Median :    0.92      Median :    0.89      Median :    0.58      Median :    0.54
## Mean   :   42.13      Mean   :   42.13      Mean   :   42.13      Mean   :   42.13
## 3rd Qu.:   21.91      3rd Qu.:   19.92      3rd Qu.:   12.27      3rd Qu.:   12.28
## Max.   :49191.15      Max.   :55692.09      Max.   :111850.87      Max.   :108726.08
##      GSM1480297      GSM1480298      GSM1480299
## Min.   :    0.000      Min.   :    0.000      Min.   :    0.000
## 1st Qu.:    0.000      1st Qu.:    0.000      1st Qu.:    0.000
## Median :    2.158      Median :    2.254      Median :    1.854
## Mean   :   42.132      Mean   :   42.132      Mean   :   42.132
## 3rd Qu.:   27.414      3rd Qu.:   26.450      3rd Qu.:   24.860
## Max.   :10489.311      Max.   :10662.486      Max.   :15194.048
##      GSM1480300      GSM1480301      GSM1480302
## Min.   :    0.000      Min.   :    0.000      Min.   :    0.000
## 1st Qu.:    0.000      1st Qu.:    0.000      1st Qu.:    0.000
## Median :    1.816      Median :    1.629      Median :    1.749
## Mean   :   42.132      Mean   :   42.132      Mean   :   42.132
## 3rd Qu.:   23.443      3rd Qu.:   23.443      3rd Qu.:   24.818
## Max.   :17434.935      Max.   :19152.728      Max.   :15997.193
```

Formatting the data

Converting from wide to long format

We will first convert the data from wide format into long format to make it easier to work with and plot with ggplot. We want just one column containing all the expression values instead of multiple columns with counts for each sample, as shown in the image below.



We can use `pivot_longer()` to easily change the format into long format.

```
seqdata <- pivot_longer(counts, cols = starts_with("GSM"), names_to = "Sample",
                        values_to = "Count")
```

We use `cols = starts_with("GSM")` to tell the function we want to reformat the columns whose names start with “GSM”. `pivot_longer()` will then reformat the specified columns into two new columns, which we’re naming “Sample” and “Count”. The `names_to = "Sample"` specifies that we want the new column containing the columns we specified with `cols` to be named “Sample”, and the `values_to = "Count"` specifies that we want the new column containing the values to be named “Count”.

We could also specify a column range to reformat. The command below would give us the same result as the previous command.

```
seqdata <- pivot_longer(counts, cols = GSM1480291:GSM1480302,
                        names_to = "Sample", values_to = "Count")
```

Alternatively, we could specify the columns we *don't* want to reformat and `pivot_longer()` will reformat all the other columns. To do that we put a minus sign “-” in front of the column names that we don’t want to reformat. This is a pretty common way to use `pivot_longer()` as sometimes it is easier to exclude columns we don’t want than include columns we do. The command below would give us the same result as the previous command.

```
seqdata <- pivot_longer(counts, cols = -c("gene_id", "gene_symbol"),
                        names_to = "Sample", values_to = "Count")
```

Here we see the function `c()` for the first time. We use this function extremely often in R when we have multiple items that we are *combining*. We will see it again in this tutorial.

Let’s have a look at the data.

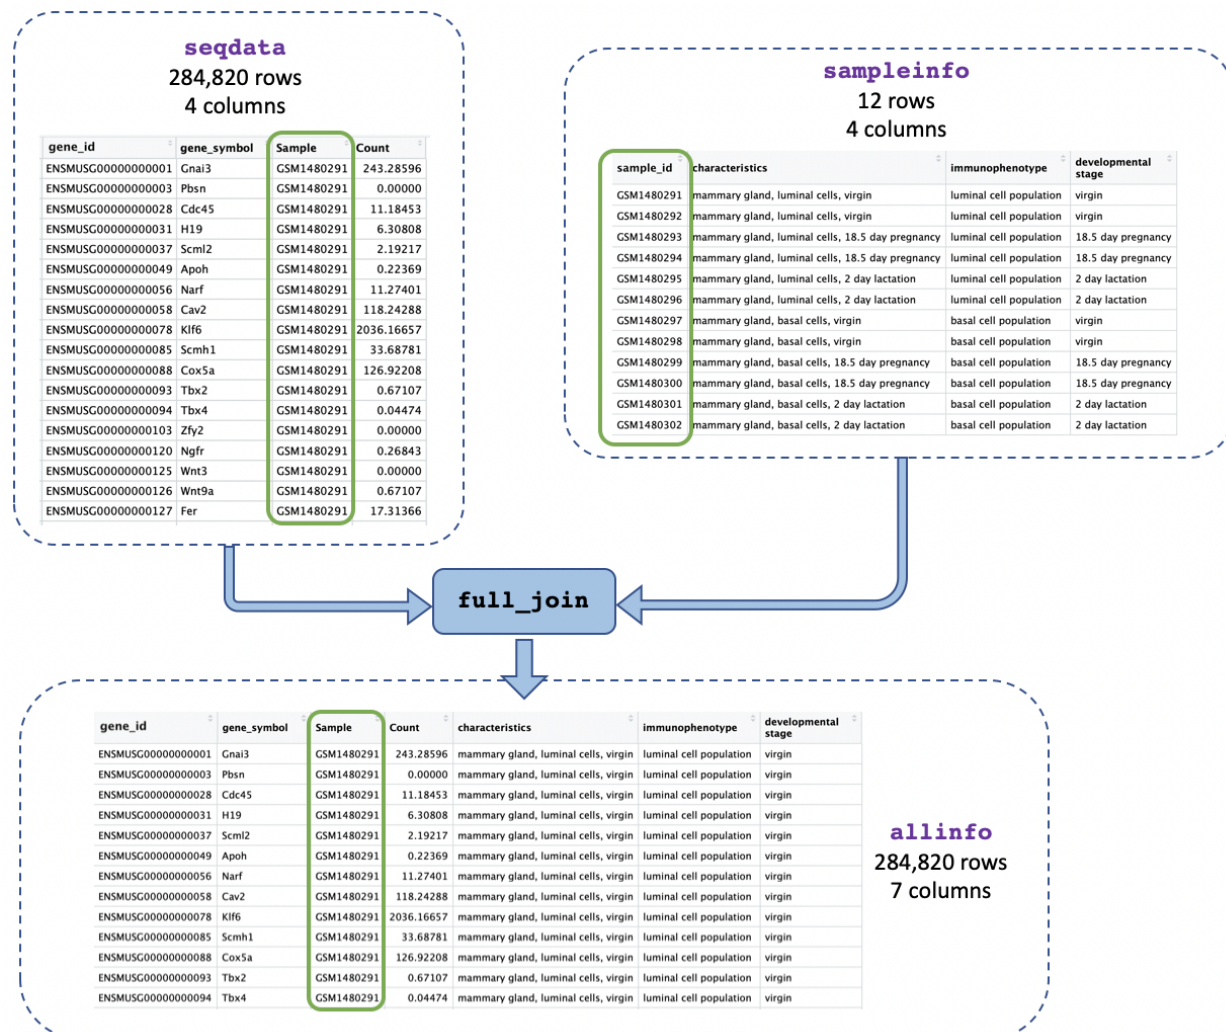
```
seqdata
```

```
## # A tibble: 284,820 x 4
##   gene_id      gene_symbol Sample      Count
```

```
##      <chr>                <chr>                <chr>      <dbl>
## 1 ENSMUSG00000000001 Gnai3      GSM1480291 243.
## 2 ENSMUSG00000000001 Gnai3      GSM1480292 256.
## 3 ENSMUSG00000000001 Gnai3      GSM1480293 240.
## 4 ENSMUSG00000000001 Gnai3      GSM1480294 217.
## 5 ENSMUSG00000000001 Gnai3      GSM1480295  84.7
## 6 ENSMUSG00000000001 Gnai3      GSM1480296  84.6
## 7 ENSMUSG00000000001 Gnai3      GSM1480297 175.
## 8 ENSMUSG00000000001 Gnai3      GSM1480298 187.
## 9 ENSMUSG00000000001 Gnai3      GSM1480299 177.
## 10 ENSMUSG00000000001 Gnai3     GSM1480300 169.
## # ... with 284,810 more rows
```

Joining two tables

Now that we've got just one column containing sample ids in both our counts and metadata objects we can join them together using the sample ids. This will make it easier to identify the categories for each sample (e.g. if it's basal cell type) and to use that information in our plots.



We will use the function `full_join()` and give it the two tables we want to join. We add `by = join_by(Sample == sample_id)` to say we want to join on the column called "Sample" in the first table

(seqdata) and the column called “sample_id” in the second table (sampleinfo) when the values match.

```
allinfo <- full_join(seqdata, sampleinfo, by = join_by(Sample == sample_id))
```

Let’s have a look at the data.

```
allinfo

## # A tibble: 284,820 x 7
##   gene_id      gene_symbol Sample      Count characteris~1 immun~2 devel~3
##   <chr>        <chr>      <chr>      <dbl> <chr>        <chr>  <chr>
## 1 ENSMUSG00000000001 Gnai3      GSM1480291 243.  mammary glan~ lumina~ virgin
## 2 ENSMUSG00000000001 Gnai3      GSM1480292 256.  mammary glan~ lumina~ virgin
## 3 ENSMUSG00000000001 Gnai3      GSM1480293 240.  mammary glan~ lumina~ 18.5 d~
## 4 ENSMUSG00000000001 Gnai3      GSM1480294 217.  mammary glan~ lumina~ 18.5 d~
## 5 ENSMUSG00000000001 Gnai3      GSM1480295  84.7 mammary glan~ lumina~ 2 day ~
## 6 ENSMUSG00000000001 Gnai3      GSM1480296  84.6 mammary glan~ lumina~ 2 day ~
## 7 ENSMUSG00000000001 Gnai3      GSM1480297 175.  mammary glan~ basal ~ virgin
## 8 ENSMUSG00000000001 Gnai3      GSM1480298 187.  mammary glan~ basal ~ virgin
## 9 ENSMUSG00000000001 Gnai3      GSM1480299 177.  mammary glan~ basal ~ 18.5 d~
## 10 ENSMUSG00000000001 Gnai3      GSM1480300 169.  mammary glan~ basal ~ 18.5 d~
## # ... with 284,810 more rows, and abbreviated variable names
## #   1: characteristics, 2: immunophenotype, 3: `developmental stage`
```

The two tables have been joined.

Exercise

1. View the help page of the `head` function and find the “Arguments” heading. What does the `n` argument do? How many rows and columns do you get with `head(sampleinfo, n = 8)`?
2. Store the output of the first 20 lines of the `counts` object in a new variable named `subset_counts`. What is the `gene_symbol` in row 20?
3. View the values in the `GSM1480291` column from your `subset_counts` object using the `$` subsetting method. What is the 5th value?
4. View the help page of the `mean` function. What is the mean of the column of values you obtained from the previous question?
5. How can you use `pivot_longer` to transform `dat` into a ‘tidy’ data called `dat_long` that contains 3 columns (sample, experiment, count).

```
dat <- tibble(sample = 1:10,
              experiment_1 = rnorm(10),
              experiment_2 = rnorm(10))
```

`dat_long` should look similar to what you get if you paste this into the console (the values in the count column will be different):

```
dat_long <- tibble(sample = rep(1:10, each=2),
                  experiment = rep(c("experiment_1", "experiment_2"), 10),
                  count = rnorm(20))
```

`dat_long`

6. Say you have another table containing sample information with samples assigned to either “Mutant” or “Control” such as:

```
sample_info <- tibble(sample = 1:100,
```



```
group = c(rep("Mutant", 50), rep("Control", 50)))
```

Join `dat_long` to `sample_info` using the common column called “sample”. How many rows do you get if you use i) `full_join`, ii) `left_join`, iii) `right_join`, iv) `inner_join`?

Plotting with ggplot2

ggplot2 is a plotting package that makes it simple to create complex plots. One really great benefit of **ggplot2** versus the older base R plotting is that we only need to make minimal changes if the underlying data change or if we decide to change our plot type, for example, from a box plot to a violin plot. This helps in creating publication quality plots with minimal amounts of adjustments and tweaking.

ggplot2 likes data in the ‘long’ format, i.e., a column for every variable, and a row for every observation, similar to what we created with `pivot_longer()`. Well-structured data will save you lots of time when making figures with **ggplot2**.

As we shall see, **ggplot** graphics are built step by step by adding new elements using the `+`. Adding layers in this fashion allows for extensive flexibility and customization of plots.

To build a **ggplot**, we use the following basic template that can be used for different types of plots. Three things are required for a **ggplot**:

```
ggplot(data= 1 , mapping=aes( 2 )) + geom_ 3 ()
```

1. The data
2. The columns in the data we want to map to visual properties (called aesthetics or `aes` in **ggplot2**)
e.g. the columns for x values, y values and colours
3. The type of plot (the `geom_`)


There are different geoms we can use to create different types of plot e.g. `geom_line()` `geom_point()`, `geom_boxplot()`. To see the geoms available take a look at the **ggplot2** help or the handy **ggplot2** cheatsheet. Or if you type “geom” in RStudio, RStudio will show you the different types of geoms you can use.

Creating a boxplot

We can make boxplots to visualise the distribution of the counts for each sample. This helps us to compare the samples and check if any look unusual.

*Note: with **ggplot** the “+” must go at the end of the line, it can’t go at the beginning.*

```
ggplot(data = allinfo, mapping = aes(x = Sample, y = Count)) +  
  geom_boxplot()
```





We have generated our first plot!

But it looks a bit weird. It's because we have some genes with extremely high counts. To make it easier to visualise the distributions we usually plot the logarithm of RNA-seq counts. We'll plot the Sample on the X axis and $\log_2(\text{Count})$ on the y axis. We can log the Counts within the `aes()`. The sample labels are also overlapping each other, we will show how to fix this later.


```
ggplot(data = allinfo, mapping = aes(x = Sample, y = log2(Count))) +
  geom_boxplot()
```

```
## Warning: Removed 84054 rows containing non-finite values (`stat_boxplot()`).
```



We get a warning here about rows containing non-finite values being removed. This is because some of the genes have a count of zero in the samples and a log of zero is undefined. We can add a small number to every count to avoid the zeros being dropped.

```
ggplot(data = allinfo, mapping = aes(x = Sample, y = log2(Count + 1))) +
  geom_boxplot()
```



The box plots show that the distributions of the samples are not identical but they are not very different.

Box plots are useful summaries, but hide the shape of the distribution. For example, if the distribution is bimodal, we would not see it in a boxplot. An alternative to the boxplot is the **violin plot**, where the shape (of the density of points) is drawn. See here for an example of how differences in distribution may be hidden in box plots but revealed with violin plots.


Exercise You can easily make different types of plots with ggplot by using different geoms. Using the same data (same x and y values), try editing the code above to make a violin plot (Hint: there's a `geom_violin`)

Colouring by categories

What if we would like to add some colour to the plot, for example, a different colour for each sample.


If we look at the `geom_boxplot` help we can see under the heading called “Aesthetics” that there's an option for colour. Let's try adding that to our plot. We'll specify we want to map the Sample column to `colour =`. As we are mapping colour to a column in our data we need to put this inside the `aes()`.

```
ggplot(data = allinfo, mapping = aes(x = Sample, y = log2(Count + 1), colour = Sample)) +
  geom_boxplot()
```



Colouring the edges wasn't quite what we had in mind. Look at the help for `geom_boxplot` to see what other aesthetic we could use. Let's try `fill =` instead.

```
ggplot(data = allinfo, mapping = aes(x = Sample, y = log2(Count + 1), fill = Sample)) +
  geom_boxplot()
```



That looks better. `fill` = is used to **fill** in areas in ggplot2 plots, whereas `colour` = is used to colour lines and points.

A really nice feature about ggplot is that we can easily colour by another variable by simply changing the column we give to `fill`.

Exercise Modify the plot above. Colour by other variables (columns) in the metadata file:

1. characteristics
2. immunophenotype
3. 'developmental stage' (As there is a space in the column name we need to use backticks around the name (")).

Note: backticks are not single quotes (''). The backtick key is usually at the top left corner of a laptop keyboard under the ESC key. Check what happens if you don't use backticks.)

Optional exercise The `geom_boxplot` function can also take in additional arguments. For example, you can decrease the size of the outlier points by using the `outlier.size` argument like so: `geom_boxplot(outlier.size = 0.5)`. View the help page for `geom_boxplot`. Can you find a way to hide outliers altogether? Plot a boxplot with hidden outliers.

Creating subplots for each gene

With ggplot we can easily make subplots using *faceting*. For example we can make stripcharts. These are a type of scatterplot and are useful when there are a small number of samples (when there are not too many points to visualise). Here we will make stripcharts plotting expression by the groups (basal virgin, basal pregnant, basal lactating, luminal virgin, luminal pregnant, luminal lactating) for each gene.

Make shorter category names

First we'll use `mutate()` to add a column with shorter group names to use in the plot, as the group names in the characteristics column are quite long.

```
allinfo <- mutate(allinfo, Group = case_when(
  str_detect(characteristics, "basal.*virgin") ~ "bvirg",
  str_detect(characteristics, "basal.*preg") ~ "bpreg",
  str_detect(characteristics, "basal.*lact") ~ "blact",
  str_detect(characteristics, "luminal.*virgin") ~ "lvirg",
  str_detect(characteristics, "luminal.*preg") ~ "lpreg",
  str_detect(characteristics, "luminal.*lact") ~ "llact"
))
```

Have a look at this data using `head()`. You should see a new column called `Group` has been added to the end.

```
head(allinfo)
```

```
## # A tibble: 6 x 8
##   gene_id      gene_symbol Sample      Count charac~1 immun~2 devel~3 Group
##   <chr>        <chr>      <chr>      <dbl> <chr>    <chr>    <chr>    <chr>
## 1 ENSMUSG00000000001 Gnai3      GSM1480291 243.   mammary~ lumina~ virgin  lvirg
## 2 ENSMUSG00000000001 Gnai3      GSM1480292 256.   mammary~ lumina~ virgin  lvirg
## 3 ENSMUSG00000000001 Gnai3      GSM1480293 240.   mammary~ lumina~ 18.5 d~ lpreg
## 4 ENSMUSG00000000001 Gnai3      GSM1480294 217.   mammary~ lumina~ 18.5 d~ lpreg
## 5 ENSMUSG00000000001 Gnai3      GSM1480295  84.7 mammary~ lumina~ 2 day ~ llact
## 6 ENSMUSG00000000001 Gnai3      GSM1480296  84.6 mammary~ lumina~ 2 day ~ llact
## # ... with abbreviated variable names 1: characteristics, 2: immunophenotype,
## #   3: `developmental stage`
```

Filter for genes of interest

We can make plots for a set of given genes.

```
mygenes <- c("Csn1s2a", "Csn1s1", "Csn2", "Glycam1", "COX1", "Trf", "Wap", "Eef1a1")
```

We filter our data for just these genes of interest. We use `%in%` to check if a value is in a set of values.

```
mygenes_counts <- filter(allinfo, gene_symbol %in% mygenes)
```

[INFO] An additional note about which genes we've chosen


The genes we've picked are the 8 genes with the highest counts summed across all samples. The code for how to get the gene symbols for these 8 genes is shown below. This code uses pipes (`%>%`) to string a series of function calls together (which is beyond the scope of this tutorial, but totally worth learning about independently!).

```
mygenes <- allinfo %>%
  group_by(gene_symbol) %>%
  summarise(Total_count = sum(Count)) %>%
  arrange(desc(Total_count)) %>%
  head(n = 8) %>%
  pull(gene_symbol)
```

Create plots for each gene


We can make boxplots for just these genes. We *facet* on the `gene_symbol` column using `facet_wrap()`. We add the tilde symbol `~` in front of the column we want to facet on.

```
ggplot(data = mygenes_counts,
       mapping = aes(x = Group, y = log2(Count + 1), fill = Group)) +
  geom_boxplot() +
  facet_wrap(~ gene_symbol)
```




The boxplots don't look good as we only have two values per group. We could just plot the individual points instead. We could use `geom_point()` to make a scatterplot.

```
ggplot(data = mygenes_counts, mapping = aes(x = Group, y = log2(Count + 1))) +
  geom_point() +
  facet_wrap(~ gene_symbol)
```




The points are overlapping so we will make a jitter plot using `geom_jitter()`. A **jitter plot** is similar to a scatter plot. It adds a small amount of random variation to the location of each point so they don't overlap. It is also quite common to combine jitter plots with other types of plot, for example, jitter with boxplot.

```
ggplot(data = mygenes_counts, mapping = aes(x = Group, y = log2(Count + 1))) +
  geom_jitter() +
  facet_wrap(~ gene_symbol)
```



We can colour the groups similar to before using `colour =`.

```
ggplot(data = mygenes_counts,
       mapping = aes(x = Group, y = log2(Count + 1), colour = Group)) +
  geom_jitter() +
  facet_wrap(~ gene_symbol)
```



Customising the plot

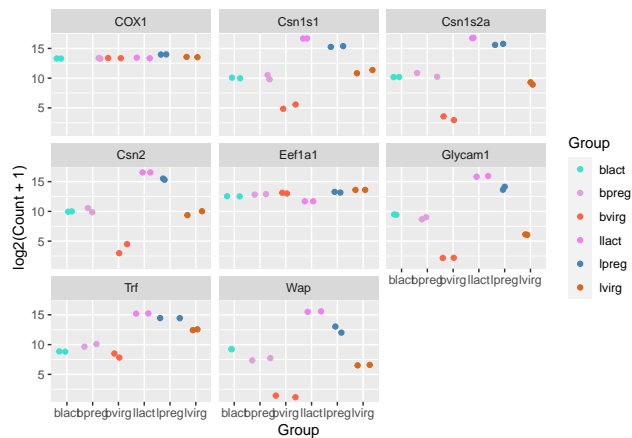
Specifying colours

We might want to change the colours. To see what colour names are available you can type `colours()`. There is also an R colours cheatsheet that shows what the colours look like.

```
mycolours <- c("turquoise", "plum", "tomato", "violet", "steelblue", "chocolate")
```

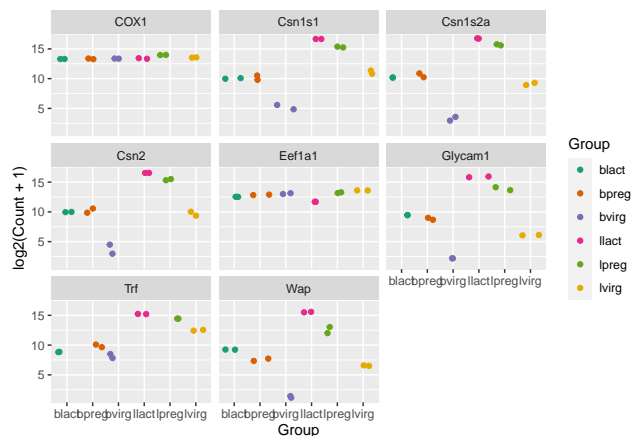
Then we then add these colours to the plot using a `+` and `scale_colour_manual(values = mycolours)`.

```
ggplot(data = mygenes_counts,
       mapping = aes(x = Group, y = log2(Count + 1), colour = Group)) +
  geom_jitter() +
  facet_wrap(~ gene_symbol) +
  scale_colour_manual(values = mycolours)
```



There are built-in colour palettes that can be handy to use, where the sets of colours are predefined. `scale_colour_brewer()` is a popular one (there is also `scale_fill_brewer()`). You can take a look at the help for `scale_colour_brewer()` to see what palettes are available. The R colours cheatsheet also shows what the colours of the palettes look like. There's one called "Dark2", let's have a look at that.


```
ggplot(data = mygenes_counts,
       mapping = aes(x = Group, y = log2(Count + 1), colour = Group)) +
  geom_jitter() +
  facet_wrap(~ gene_symbol) +
  scale_colour_brewer(palette = "Dark2")
```



Axis labels and Title

We can change the axis labels and add a title with `labs()`. To change the x axis label we use `labs(x = "New name")`. To change the y axis label we use `labs(y = "New name")` or we can change them all at the same time.


```
ggplot(data = mygenes_counts,
       mapping = aes(x = Group, y = log2(Count + 1), colour = Group)) +
  geom_jitter() +
  facet_wrap(~ gene_symbol) +
  labs(x = "Cell type and stage", y = "Count", title = "Mammary gland RNA-seq data")
```



Themes

We can adjust the text on the x axis (the group labels) by turning them 90 degrees so we can read the labels better. To do this we modify the ggplot theme. Themes are the non-data parts of the plot.


```
ggplot(data = mygenes_counts,
       mapping = aes(x = Group, y = log2(Count + 1), colour = Group)) +
  geom_jitter() +
  facet_wrap(~ gene_symbol) +
  labs(x = "Cell type and stage", y = "Count", title = "Mammary gland RNA-seq data") +
  theme(axis.text.x = element_text(angle = 90))
```




We can remove the grey background and grid lines.

There are also a lot of built-in themes. Let's have a look at a couple of the more widely used themes. The default ggplot theme is `theme_grey()`.

```
ggplot(data = mygenes_counts,
       mapping = aes(x = Group, y = log2(Count + 1), colour = Group)) +
  geom_jitter() +
  facet_wrap(~ gene_symbol) +
  labs(x = "Cell type and stage", y = "Count", title = "Mammary gland RNA-seq data") +
  theme_bw() +
  theme(axis.text.x = element_text(angle = 90))
```


```
ggplot(data = mygenes_counts,
       mapping = aes(x = Group, y = log2(Count + 1), colour = Group)) +
  geom_jitter() +
  facet_wrap(~ gene_symbol) +
  labs(x = "Cell type and stage", y = "Count", title = "Mammary gland RNA-seq data") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 90))
```



There are many themes available, you can see some in the R graph gallery.

We can also modify parts of the theme individually. We can remove the grey background and grid lines with the code below.

```
ggplot(data = mygenes_counts,
       mapping = aes(x = Group, y = log2(Count + 1), colour = Group)) +
  geom_jitter() +
  facet_wrap(~ gene_symbol) +
  labs(x = "Cell type and stage", y = "Count", title = "Mammary gland RNA-seq data") +
  theme(axis.text.x = element_text(angle = 90)) +
  theme(panel.background = element_blank(),
        panel.grid.major = element_blank(),
        panel.grid.minor = element_blank())
```



Order of categories

The groups have been plotted in alphabetical order on the x axis and in the legend (that is the default order), however, we may want to change the order. We may prefer to plot the groups in order of stage, for example, basal virgin, basal pregnant, basal lactate, luminal virgin, luminal pregnant, luminal lactate.

First let's make an object with the group order that we want.

```
group_order <- c("bvirg", "bpreg", "blact", "lvirg", "lpreg", "llact")
```

Next we need to make a column with the groups into an R data type called a **factor**. Factors in R are a special data type used to specify categories, you can read more about them in the R for Data Science book. The names of the categories are called the factor **levels**.

We'll add another column called "Group_f" where we'll make the Group column into a factor and specify what order we want the levels of the factor.

```
mygenes_counts <- mutate(mygenes_counts, Group_f = factor(Group, levels = group_order))
```

Take a look at the data. As the table is quite wide we can use `select()` to select just the columns we want to view.

```
select(mygenes_counts, gene_id, Group, Group_f)
```

```
## # A tibble: 96 x 3
##   gene_id          Group Group_f
##   <chr>          <chr> <fct>
## 1 ENSMUSG00000000381 lvirg lvirg
## 2 ENSMUSG00000000381 lvirg lvirg
## 3 ENSMUSG00000000381 lpreg lpreg
## 4 ENSMUSG00000000381 lpreg lpreg
## 5 ENSMUSG00000000381 llact llact
## 6 ENSMUSG00000000381 llact llact
## 7 ENSMUSG00000000381 bvirg bvirg
## 8 ENSMUSG00000000381 bvirg bvirg
## 9 ENSMUSG00000000381 bpreg bpreg
## 10 ENSMUSG00000000381 bpreg bpreg
## # ... with 86 more rows
```

Notice that the Group column has `<chr>` under the heading, that indicates is a character data type, while the Group_f column has `<fct>` under the heading, indicating it is a factor data type. The `str()` command that we saw previously is useful to check the data types in objects.

```
str(mygenes_counts)
```

```
## tibble [96 x 9] (S3: tbl_df/tbl/data.frame)
## $ gene_id      : chr [1:96] "ENSMUSG00000000381" "ENSMUSG00000000381" "ENSMUSG00000000381" "E
## $ gene_symbol  : chr [1:96] "Wap" "Wap" "Wap" "Wap" ...
## $ Sample      : chr [1:96] "GSM1480291" "GSM1480292" "GSM1480293" "GSM1480294" ...
## $ Count       : num [1:96] 90.2 95.6 4140.3 8414.4 49204.9 ...
## $ characteristics : chr [1:96] "mammary gland, luminal cells, virgin" "mammary gland, luminal ce
## $ immunophenotype : chr [1:96] "luminal cell population" "luminal cell population" "luminal cell
## $ developmental stage: chr [1:96] "virgin" "virgin" "18.5 day pregnancy" "18.5 day pregnancy" ...
## $ Group       : chr [1:96] "lvirg" "lvirg" "lpreg" "lpreg" ...
## $ Group_f     : Factor w/ 6 levels "bvirg","bpreg",...: 4 4 5 5 6 6 1 1 2 2 ...
```

str() shows us Group_f column is a Factor with 6 levels (categories).


We can check the factor levels of a column as below.

```
levels(mygenes_counts$Group_f)
```

```
## [1] "bvirg" "bpreg" "blact" "lvirg" "lpreg" "llact"
```

The levels are in the order that we want, so we can now change our plot to use the “Group_f” column instead of Group column (change x = and colour =).

```
ggplot(data = mygenes_counts,
       mapping = aes(x = Group_f, y = log2(Count + 1), colour = Group_f)) +
  geom_jitter() +
  facet_wrap(~ gene_symbol) +
  labs(x = "Cell type and stage", y = "Count", title = "Mammary gland RNA-seq data") +
  theme(axis.text.x = element_text(angle = 90)) +
  theme(panel.background = element_blank(),
        panel.grid.major = element_blank(),
        panel.grid.minor = element_blank())
```



We could do similar if we wanted to have the genes in the facets in a different order. For example, we could add another column called “gene_symbol_f” where we make the gene_symbol column into a factor, specifying the order of the levels.

Exercise

1. Make a colourblind-friendly plot using the colourblind-friendly palettes here.
2. Create a plot (any plot whatsoever) and share it with the class by pasting the image in the Google

Docs link provided in your workshop. Your plot should use the `subtitle` argument in the `labs` function to add a unique identifier (e.g. a message and your name or initials) which is displayed below the title.

Tip: An easy way to copy your plot in RStudio is using the plot pane's export option and selecting "Copy to Clipboard...". You can then paste it into the provided Google document.

Saving plots

We can save plots interactively by clicking Export in the Plots window and saving as e.g. "myplot.pdf". Or we can output plots to pdf using `pdf()` followed by `dev.off()`. We put our plot code after the call to `pdf()` and before closing the plot device with `dev.off()`.

Let's save our last plot.

```
pdf("myplot.pdf")
ggplot(data = mygenes_counts,
       mapping = aes(x = Group_f, y = log2(Count + 1), colour = Group_f)) +
  geom_jitter() +
  facet_wrap(~ gene_symbol) +
  labs(x = "Cell type and stage", y = "Count", title = "Mammary gland RNA-seq data") +
  theme(axis.text.x = element_text(angle = 90)) +
  theme(panel.background = element_blank(),
        panel.grid.major = element_blank(),
        panel.grid.minor = element_blank())
dev.off()
```

Exercise

1. Download the raw counts for this dataset from GREIN
 - a. Make a boxplot. Do the samples look any different to the normalised counts?
 - b. Make subplots for the same set of 8 genes. Do they look any different to the normalised counts?
2. Download the normalised counts for the GSE63310 dataset from GREIN. Make boxplots colouring the samples using different columns in the metadata file.

Session Info

The last thing we'll do run the `sessionInfo()` function. This function prints out details about your working environment such as the version of R you're running, loaded packages, and package versions. Printing out `sessionInfo()` at the end of your analysis is good practice as it helps with reproducibility in the future.

```
sessionInfo()

## R version 4.2.1 (2022-06-23)
## Platform: x86_64-apple-darwin17.0 (64-bit)
## Running under: macOS Big Sur ... 10.16
##
## Matrix products: default
## BLAS:   /Library/Frameworks/R.framework/Versions/4.2/Resources/lib/libRblas.0.dylib
## LAPACK: /Library/Frameworks/R.framework/Versions/4.2/Resources/lib/libRlapack.dylib
##
## locale:
##  [1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8
##
## attached base packages:
```

```
## [1] stats      graphics  grDevices utils      datasets  methods  base
##
## other attached packages:
## [1] forcats_0.5.2  stringr_1.5.0  dplyr_1.1.0    purrr_1.0.1
## [5] readr_2.1.3    tidyr_1.3.0    tibble_3.1.8   ggplot2_3.4.0
## [9] tidyverse_1.3.2
##
## loaded via a namespace (and not attached):
## [1] lubridate_1.9.0  assertthat_0.2.1  digest_0.6.30
## [4] utf8_1.2.2      R6_2.5.1          cellranger_1.1.0
## [7] backports_1.4.1  reprex_2.0.2      evaluate_0.18
## [10] httr_1.4.4       pillar_1.8.1      rlang_1.0.6
## [13] googlesheets4_1.0.1 readxl_1.4.1      rstudioapi_0.14
## [16] rmarkdown_2.18   labeling_0.4.2    googledrive_2.0.0
## [19] bit_4.0.5        munsell_0.5.0     broom_1.0.1
## [22] compiler_4.2.1   modelr_0.1.10     xfun_0.35
## [25] pkgconfig_2.0.3  htmltools_0.5.3   tidyselect_1.2.0
## [28] fansi_1.0.3      crayon_1.5.2      tzdb_0.3.0
## [31] dbplyr_2.2.1     withr_2.5.0       grid_4.2.1
## [34] jsonlite_1.8.3   gtable_0.3.1      lifecycle_1.0.3
## [37] DBI_1.1.3        magrittr_2.0.3    scales_1.2.1
## [40] cli_3.4.1        stringi_1.7.8     vroom_1.6.0
## [43] farver_2.1.1     fs_1.5.2          xml2_1.3.3
## [46] ellipsis_0.3.2   generics_0.1.3    vctrs_0.5.2
## [49] RColorBrewer_1.1-3 tools_4.2.1       bit64_4.0.5
## [52] glue_1.6.2       hms_1.1.2         parallel_4.2.1
## [55] fastmap_1.1.0    yaml_2.3.6        timechange_0.1.1
## [58] colorspace_2.0-3 gargle_1.2.1      rvest_1.0.3
## [61] knitr_1.41       haven_2.5.1
```

Key Points

- Tabular data can be loaded into R with the tidyverse functions `read_csv()` and `read_tsv()`
- Tidyverse functions such as `pivot_longer()`, `mutate()`, `filter()`, `select()`, `full_join()` can be used to manipulate data
- A ggplot has 3 components: data (dataset), mapping (columns to plot) and geom (type of plot). Different types of plots include `geom_point()`, `geom_jitter()`, `geom_line()`, `geom_boxplot()`, `geom_violin()`.
- `facet_wrap()` can be used to make subplots of the data
- The aesthetics of a ggplot can be modified, such as colouring by different columns in the dataset, adding labels or changing the background

Further Reading

A short intro to R and tidyverse
 Top 50 Ggplot Visualisations
 R for Data Science